

Adit Mohindra

aditmohindra.com | linkedin.com/in/adit-mohindra | github.com/aditmohindra | San Francisco, CA

Machine Learning Engineer with experience building production ML systems, ETL pipelines, and executive-facing analytics.

EDUCATION

University of California, Irvine

B.S. Data Science and Statistics — Dean's List | 1st Place, ICS Project Expo

EXPERIENCE

Cotality

Jan. 2025 – Present

Machine Learning Engineer - Professional, Business Analysis

- **Built a production ensemble forecasting system for a \$4M operational planning platform, delivering the first model improvement**, reducing WAPE 6.5% across 1,000+ independent customer forecasts and 1M+ monthly units
- **Reduced forecast error by 37% MAE on high-variance customer segments while preserving legacy strength** by developing a worktype-level model router through historical backtesting and live regime monitoring
- **Owned multiple forecasting platforms end to end**, including data ingestion, validation, model evaluation, and executive-facing Tableau dashboards reporting forecast accuracy and business impact
- **Built MLOps infrastructure using ZenML** to automate data validation, model retraining, monthly customer routing, monitoring, and production fallback workflows
- **Engineered automated Python ETL pipelines replacing legacy Excel/VBScript processes**, improving vendor data processing time by 82% and eliminating downstream reporting errors across critical workflows

Machine Learning Researcher

- **Awarded 1st Place at UC Irvine's ICS Project Expo** for designing an end-to-end ML product forecasting ZIP-level real estate market trends and investment attractiveness using proprietary Cotality data.
- **Built a modular ML pipeline and interactive Streamlit application**, enabling stakeholders to explore temporal and regional trends, visualize forecasts, and customize market-hotness weighting in real time

Procon Analytics

Jul. 2024 – Aug. 2025

Data Science Intern

- **Developed a customer churn-risk model, identifying \$2M+ in revenue at risk** and enabled targeted retention strategies across product lines and customer segments
- **Unified 100K+ fragmented customer records through Python ETL and SQL fuzzy matching**, improving enterprise data quality 42% and reducing reporting inconsistencies 19%.
- **Deployed real-time Tableau dashboards** enabling sales and customer experience teams to explore churn-risk trends across products and regions

PROJECTS

Spotify Atlas

[Project Link](#)

- **Built a music intelligence ML product that turns over a decade of user Spotify history into an explorable galaxy** of interconnected music, using semantic embeddings and high-dimensional clustering to map relationships between tracks and form communities around shared music culture and identity
- **Evaluated competing embedding and clustering configurations** using silhouette score, coverage, coherence scoring, and track-level inspection, selecting production runs that balanced statistical separation with human interpretability

TECHNICAL SKILLS

- **Languages & Analytics:** Python, SQL, R, Excel, Git, Jupyter, NumPy, Pandas, TypeScript, Next.js
- **ML & Forecasting:** Time-Series Forecasting, Backtesting, Ensemble Methods, XGBoost, Random Forest, SARIMA, PyMC, Chronos, LSTM, PyTorch, Scikit-learn
- **MLOps & Data Engineering:** MLflow, ETL, Data Validation, ZenML, Prefect, Docker, Google BigQuery, PostgreSQL, MySQL, AWS S3, EC2, SageMaker, Airflow, Athena, Snowflake, PySpark, FastAPI, CI/CD
- **Visualization & Reporting:** Tableau, Streamlit, Plotly Dash, Matplotlib, Seaborn, Power BI